
Table of Contents

Preface.....	xiii
--------------	------

Part I. Storage Engines

1. Introduction and Overview.....	7
DBMS Architecture	8
Memory- Versus Disk-Based DBMS	10
Durability in Memory-Based Stores	11
Column- Versus Row-Oriented DBMS	12
Row-Oriented Data Layout	13
Column-Oriented Data Layout	14
Distinctions and Optimizations	15
Wide Column Stores	15
Data Files and Index Files	17
Data Files	18
Index Files	18
Primary Index as an Indirection	20
Buffering, Immutability, and Ordering	21
Summary	22
2. B-Tree Basics.....	25
Binary Search Trees	26
Tree Balancing	27
Trees for Disk-Based Storage	28
Disk-Based Structures	29
Hard Disk Drives	30
Solid State Drives	30

On-Disk Structures	32
Ubiquitous B-Trees	33
B-Tree Hierarchy	35
Separator Keys	36
B-Tree Lookup Complexity	37
B-Tree Lookup Algorithm	38
Counting Keys	38
B-Tree Node Splits	39
B-Tree Node Merges	41
Summary	42
3. File Formats.....	45
Motivation	46
Binary Encoding	47
Primitive Types	47
Strings and Variable-Size Data	49
Bit-Packed Data: Booleans, Enums, and Flags	49
General Principles	50
Page Structure	52
Slotted Pages	52
Cell Layout	54
Combining Cells into Slotted Pages	56
Managing Variable-Size Data	57
Versioning	58
Checksumming	59
Summary	60
4. Implementing B-Trees.....	61
Page Header	61
Magic Numbers	62
Sibling Links	62
Rightmost Pointers	63
Node High Keys	64
Overflow Pages	65
Binary Search	67
Binary Search with Indirection Pointers	67
Propagating Splits and Merges	68
Breadcrumbs	69
Rebalancing	70
Right-Only Appends	71
Bulk Loading	72
Compression	73

Vacuum and Maintenance	74
Fragmentation Caused by Updates and Deletes	75
Page Defragmentation	76
Summary	76
5. Transaction Processing and Recovery.....	79
Buffer Management	81
Caching Semantics	83
Cache Eviction	83
Locking Pages in Cache	84
Page Replacement	85
Recovery	88
Log Semantics	90
Operation Versus Data Log	91
Steal and Force Policies	91
ARIES	92
Concurrency Control	93
Serializability	94
Transaction Isolation	95
Read and Write Anomalies	95
Isolation Levels	96
Optimistic Concurrency Control	98
Multiversion Concurrency Control	99
Pessimistic Concurrency Control	99
Lock-Based Concurrency Control	100
Summary	108
6. B-Tree Variants.....	111
Copy-on-Write	112
Implementing Copy-on-Write: LMDB	113
Abstracting Node Updates	113
Lazy B-Trees	114
WiredTiger	114
Lazy-Adaptive Tree	116
FD-Trees	117
Fractional Cascading	118
Logarithmic Runs	119
Bw-Trees	120
Update Chains	121
Taming Concurrency with Compare-and-Swap	121
Structural Modification Operations	122
Consolidation and Garbage Collection	123

Cache-Oblivious B-Trees	124
van Emde Boas Layout	125
Summary	127
7. Log-Structured Storage.....	129
LSM Trees	130
LSM Tree Structure	132
Updates and Deletes	136
LSM Tree Lookups	137
Merge-Iteration	137
Reconciliation	140
Maintenance in LSM Trees	141
Read, Write, and Space Amplification	143
RUM Conjecture	144
Implementation Details	145
Sorted String Tables	145
Bloom Filters	146
SkipList	148
Disk Access	150
Compression	151
Unordered LSM Storage	152
Bitcask	153
WiscKey	154
Concurrency in LSM Trees	155
Log Stacking	157
Flash Translation Layer	157
Filesystem Logging	159
LLAMA and Mindful Stacking	160
Open-Channel SSDs	161
Summary	162
Part I Conclusion.....	165

Part II. Distributed Systems

8. Introduction and Overview.....	171
Concurrent Execution	171
Shared State in a Distributed System	173
Fallacies of Distributed Computing	174
Processing	175
Clocks and Time	176

State Consistency	177
Local and Remote Execution	178
Need to Handle Failures	178
Network Partitions and Partial Failures	179
Cascading Failures	180
Distributed Systems Abstractions	181
Links	182
Two Generals' Problem	187
FLP Impossibility	189
System Synchrony	190
Failure Models	191
Crash Faults	191
Omission Faults	192
Arbitrary Faults	193
Handling Failures	193
Summary	193
9. Failure Detection.....	195
Heartbeats and Pings	196
Timeout-Free Failure Detector	197
Outsourced Heartbeats	198
Phi-Accrual Failure Detector	199
Gossip and Failure Detection	200
Reversing Failure Detection Problem Statement	201
Summary	202
10. Leader Election.....	205
Bully Algorithm	207
Next-In-Line Failover	208
Candidate/Oldinary Optimization	209
Invitation Algorithm	210
Ring Algorithm	211
Summary	212
11. Replication and Consistency.....	215
Achieving Availability	216
Infamous CAP	216
Use CAP Carefully	217
Harvest and Yield	218
Shared Memory	219
Ordering	221
Consistency Models	222

Strict Consistency	223
Linearizability	223
Sequential Consistency	227
Causal Consistency	229
Session Models	233
Eventual Consistency	234
Tunable Consistency	235
Witness Replicas	236
Strong Eventual Consistency and CRDTs	238
Summary	240
12. Anti-Entropy and Dissemination.....	243
Read Repair	245
Digest Reads	246
Hinted Handoff	246
Merkle Trees	247
Bitmap Version Vectors	248
Gossip Dissemination	250
Gossip Mechanics	251
Overlay Networks	251
Hybrid Gossip	253
Partial Views	254
Summary	255
13. Distributed Transactions.....	257
Making Operations Appear Atomic	258
Two-Phase Commit	259
Cohort Failures in 2PC	261
Coordinator Failures in 2PC	262
Three-Phase Commit	264
Coordinator Failures in 3PC	265
Distributed Transactions with Calvin	266
Distributed Transactions with Spanner	268
Database Partitioning	270
Consistent Hashing	271
Distributed Transactions with Percolator	272
Coordination Avoidance	275
Summary	277
14. Consensus.....	279
Broadcast	280
Atomic Broadcast	281

Virtual Synchrony	282
Zookeeper Atomic Broadcast (ZAB)	283
Paxos	285
Paxos Algorithm	286
Quorums in Paxos	287
Failure Scenarios	288
Multi-Paxos	291
Fast Paxos	292
Egalitarian Paxos	293
Flexible Paxos	296
Generalized Solution to Consensus	297
Raft	300
Leader Role in Raft	302
Failure Scenarios	304
Byzantine Consensus	305
PBFT Algorithm	306
Recovery and Checkpointing	309
Summary	309
Part II Conclusion	313
A. Bibliography	317
Index	337

